

Blockchain

Nicoleta Preda

January 8, 2019

Agenda

Crypto Principles (from Eduard Felten's tutorial)

- Hash Functions

- Hash Pointers and Data Structures

- Digital Signatures

- Public Keys as Identities

- Simple Cryptocurrencies

Distributed Architecture

Bitcoin

Conclusion

Hash Functions

A hash function:


- ▶ takes any string as input
- ▶ fixed-size output (e.g., Bitcoin uses 256 bits)
- ▶ efficiently computable

Security properties:


- ▶ collision-free
- ▶ hiding
- ▶ puzzle-friendly

Hash Property 1: Collision-Free

Nobody can find x and y such that
 $x \neq y$ and $H(x) = H(y)$



Collisions do exist ...



... but can anyone find them?

How to find a collision?

try 2^{130} randomly chosen inputs

99.8% chance that two of them will collide

This works no matter what H is ...

... but it takes too long to matter

Is there a faster way to find collisions?

- ▶ For some possible H 's, yes. (Example ?)
- ▶ For others, we don't know of one.

No H has been proven collision-free.

Application: Hash as message digest.

Useful to verify if two very large files are identical.

If we know $H(x) = H(y)$,
it's safe to assume that $x = y$.

To recognise a file that we saw before,
just remember its hash.

Useful because the hash is small.

Hash Property 2: Hiding

We want something like this:


Given $H(x)$, it is infeasible to find x .

Hash Property 2: Hiding

We want something like this:

Given $H(x)$, it is infeasible to find x .

Exercise:



How easy is to find x here?

Hiding Property:

If r is chosen from a probability distribution that has high min-entropy, then given $H(r \mid x)$, it is infeasible to find x .

High min-entropy means that the distribution is “very spread out”, so that no particular value is chosen with more than negligible probability.

Application: Commitment

- ▶ Want to “seal a value in an envelope”, and “open the envelope” later.
- ▶ Commit to a value, reveal it later.

Commitment API:

- To seal msg in envelope:
select key, a random 256-bit value,
 $\text{commit}(\text{msg}) \rightarrow (\text{com} = \text{H}(\text{key} \mid \text{msg}), \text{H}(\text{key}))$, then publish com
- To open envelope:
publish key, msg
anyone can check validity: $\text{H}(\text{key} \mid \text{msg}) == \text{com}$

Security Properties for the Commitment API

Hiding: Given the result of $H(\text{key} \mid \text{msg})$, infeasible to find msg .

Binding: Infeasible to find $\text{msg} \neq \text{msg}'$ such that
 $H(\text{key} \mid \text{msg}) == H(\text{key} \mid \text{msg}')$

Hash Property 3: Puzzle-Friendly

Puzzle-friendly:

For every possible output value y , if k is chosen from a distribution with high min-entropy, then it is infeasible to find x such that $H(k \parallel x) = y$.

Application: Search Puzzle

Puzzle-friendly property implies that no solving strategy is much better than trying random values of x .

Given a “puzzle ID” id (from high min-entropy distrib.),
and a target set Y :

Try to find a “solution” x such that
 $H(id \parallel x) \in Y$.

Agenda

Crypto Principles (from Eduard Felten's tutorial)

- Hash Functions

- Hash Pointers and Data Structures

- Digital Signatures

- Public Keys as Identities

- Simple Cryptocurrencies

Distributed Architecture

Bitcoin

Conclusion

Principle

hash pointer is:

- ▶ pointer to where some info is stored, and
- ▶ hash of the info (cryptographic hash)


if we have a hash pointer, we can

- ▶ ask to get the info back, and
- ▶ verify that it hasn't changed



Distributed Data Structure: Block Chain

linked list with hash pointers = “block chain”

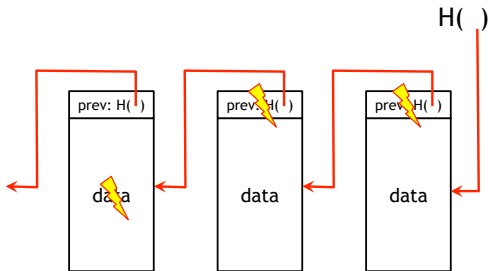


Use-case: tamper-evident log

Data is added to the end.


No one can change what has been previously added.

Detecting Tampering




Merkle Tree

binary tree with hash pointers = "Merkle tree"



Proving Membership in a Merkle Tree

show $O(\log n)$ items



Advantages of Merkle Trees

- ▶ Just need to remember the root hash
- ▶ Can verify membership in $O(\log n)$ time/space

Variant: sorted Merkle tree

- ▶ can verify non-membership in $O(\log n)$
- ▶ show items before, after the missing one

Agenda

Crypto Principles (from Eduard Felten's tutorial)

- Hash Functions

- Hash Pointers and Data Structures

- Digital Signatures

- Public Keys as Identities

- Simple Cryptocurrencies

Distributed Architecture

Bitcoin

Conclusion

Digital Signatures

- ▶ Only you can sign, but anyone can verify
- ▶ Signature is tied to a particular document
can't be cut-and-pasted to another doc

API for digital signatures

- ▶ `generateKeys(keysize) → (sk, pk)`
sk: secret signing key
pk: public verification key
- ▶ `sign(sk, message) → sig`
- ▶ `verify(pk, message, sig) → isValid`

Bitcoin uses ECDSA standard (Elliptic Curve Digital Signature Algorithm)

Requirements for Digital Signatures

“valid signatures verify”:

$\text{verify}(\text{pk}, \text{message}, \text{sign}(\text{sk}, \text{message})) == \text{true}$

“can't forge signatures”:

adversary who:

knows pk

gets to see signatures on messages of his choice

can't produce a verifiable signature on another message

Practical Constrains of Digital Signature Algorithms

- ▶ Algorithms are randomized
 - need a good source of randomness
- ▶ Have an upper bound limit on the size of the message
 - fix: sign the result of the hash function

Agenda

Crypto Principles (from Eduard Felten's tutorial)

Hash Functions

Hash Pointers and Data Structures

Digital Signatures

Public Keys as Identities

Simple Cryptocurrencies

Distributed Architecture

Bitcoin

Conclusion

public key = an identity (= address in Bitcoin)

Decentralised identity management

- ▶ to “speak for” pk, you must know matching secret key sk
- ▶ anybody can make a new identity at any time
make as many as you want!
- ▶ no central point of coordination

Privacy

- ▶ Addresses not directly connected to real-world identity.
- ▶ But observers can link together an address's activity over time, make inferences.

Agenda

Crypto Principles (from Eduard Felten's tutorial)

- Hash Functions

- Hash Pointers and Data Structures

- Digital Signatures

- Public Keys as Identities

- Simple Cryptocurrencies

Distributed Architecture

Bitcoin

Conclusion

Two Simple Cryptocurrencies

► GoofyCoin

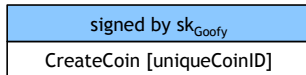


► ScroogeCoin



GoofyCoin

- ▶ Goofy can create new coins (by design).
- ▶ The new coin belongs to Goofy



GoofyCoin (continue)

- ▶ A coin's owner can spend it.
- ▶ Goofy pays Alice.
- ▶ Alice owns the coin now.

signed by sk_{Goofy}
Pay to $pk_{\text{Alice}} : H(\cdot)$




signed by sk_{Goofy}
CreateCoin [uniqueCoinID]



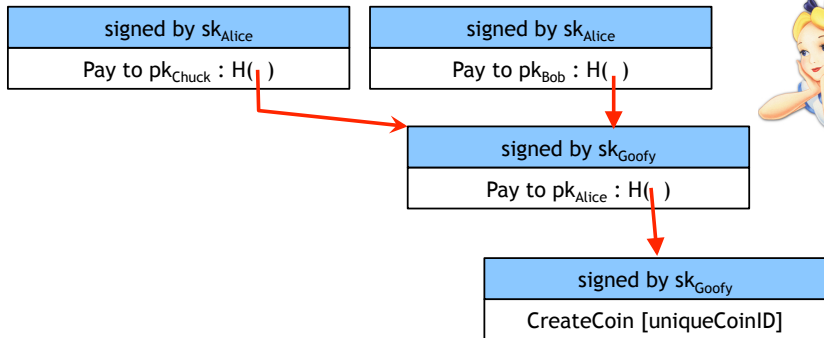
GoofyCoin (continue)

- ▶ The recipient (Alice) can pass on the coin again.
- ▶ Bob owns the coin now.



GoofyCoin Issue: Double-Spending Attack

- ▶ Alice passes the same coin to Bob and Chuck.
- ▶ Which of the two blocks is valid?



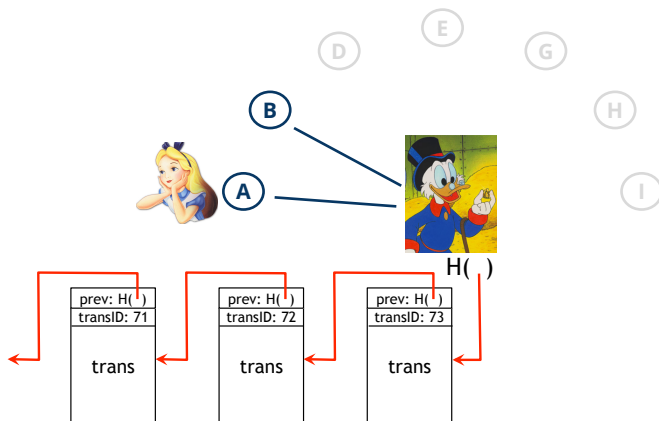
Double-Spending: One of the main design challenges in digital currencies!



- ▶ Centralised Solution
- ▶ Decentralised Solution (Bitcoin)

Scrooge Coin: A Centralised Solution

- ▶ A central bank (Scrooge) maintains a history of all transactions.
- ▶ History of all transactions: a block chain, signed by Scrooge



Detail

- ▶ We consider one transaction per block-chain.
- ▶ In reality, for optimisation, a block contains multiple transactions.

Example Transaction

- ▶ CreateCoins Transaction creates new coins.
- ▶ Is valid because Scrooge, the central authority, decides so.
- ▶ Everyone trusts Scrooge.

transID: 73 type:CreateCoins		
coins created:		
num	value	recipient
0	3.2	0x...
1	1.4	0x...
2	7.1	0x...

← coinID 73(0)

← coinID 73(1)

← coinID 73(2)

Example Transaction

- ▶ CreateCoins Transaction creates new coins.
- ▶ Is valid because Scrooge, the central authority, decides so.
- ▶ Everyone trusts Scrooge.

transID: 73 type:CreateCoins		
coins created:		
num	value	recipient
0	3.2	0x...
1	1.4	0x...
2	7.1	0x...

← coinID 73(0)

← coinID 73(1)

← coinID 73(2)

Example Transaction (continue)

- ▶ PayCoins transaction consumes (and destroys) some coins, and creates new coins of the same total value
- ▶ The transaction is valid if:
 - consumed coins valid, not already consumed,
 - total value out = total value in, and
 - signed by owners of all consumed coins

transID: 73 type:PayCoins		
consumed coinIDs: 68(1), 42(0), 72(3)		
coins created:		
num	value	recipient
0	3.2	0x...
1	1.4	0x...
2	7.1	0x...
Signatures owners of consumed		

Coins are Immutable

How to subdivide a coins?

- ▶ Create a new PayCoins transaction.
- ▶ Pay out two new coins to yourself.

- ▶ Issue with the centralised solution:
You need to trust Scrooge (the central bank).
- ▶ Solution:
Distributed, self-organising peer-to-peer solution.
No central control, no unique point of failure.

Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

- P2P Overlay Networks

- Dissemination Algorithms

- Security

- Consistency

Bitcoin

Conclusion

Peer-to-Peer (P2P) Networks

Advantages

- ▶ Robustness: the disconnection of a peer is without consequences.
- ▶ Responsibility and costs are shared among peers.


Disadvantages

- ▶ Complex coordination as peers come and leave (response time).
- ▶ Malicious participants may take control over the network (security).
- ▶ No control of the dissemination of information (privacy).

P2P Topology


Unstructured

- ▶ Random connections between peers
- ▶ No overhead for maintaining the structure



Structured: Distributed Hash Tables (DHT)


- ▶ Ring of peers (hashed ids)
- ▶ Every peer assigned specific links
- ▶ Superimposed binary look-up trees
- ▶ Cached links for fast look-up
- ▶ High overhead when peers churn



P2P Topology

Unstructured but not random

- ▶ Emergent Super-Peers
- ▶ Possible self-organization: peers connect “well connected” peers
- ▶ Some overhead caused by super-peers



Blockchain: Challenges and Approaches

- ▶ Q: Which topology is more suited for a blockchain?

Unstructured P2P: Join the Network

Challenge: no central entity to contact

Bitcoin's Approach:

- ▶ Use pre-configured IP addresses.
- ▶ Rely on volunteers to host DNS servers.
- ▶ Then, learn IP addresses of other nodes.
- ▶ A peer connects to a certain number of peers (default: 8)
- ▶ Reciprocally, a peer accepts incoming connections (on average 30).
- ▶ Periodically ping peers to check liveness
- ▶ Swap so-so nodes for better ones (several hours).

Blockchain: Challenges and Approaches

- ▶ Q: Where is the ledger stored?

Blockchain: Challenges and Approaches

- ▶ Q: Where is the ledger stored?
- ▶ A: Every peer stores the history of the transactions.

Blockchain: Challenges and Approaches

- ▶ Q: Where is the ledger stored?
- ▶ A: Every peer stores the history of the transactions.
- ▶ Q: How are new transactions handled?
- ▶ Example: "Alice wants to pay Bob 5 cents."

Blockchain: Challenges and Approaches

- ▶ Q: Where is the ledger stored?
- ▶ A: Every peer stores the history of the transactions.
- ▶ Q: How are new transactions handled?
- ▶ Example: "Alice wants to pay Bob 5 cents."
- ▶ A: A p2p algorithm is used to "broadcast" the message

Blockchain: Challenges and Approaches

- ▶ Q: Where is the ledger stored?
- ▶ A: Every peer stores the history of the transactions.
- ▶ Q: How are new transactions handled?
- ▶ Example: "Alice wants to pay Bob 5 cents."
- ▶ A: A p2p algorithm is used to "broadcast" the message
- ▶ Q: What happens when a peer receives two identical messages?

Blockchain: Challenges and Approaches

- ▶ Q: Where is the ledger stored?
- ▶ A: Every peer stores the history of the transactions.
- ▶ Q: How are new transactions handled?
- ▶ Example: "Alice wants to pay Bob 5 cents."
- ▶ A: A p2p algorithm is used to "broadcast" the message
- ▶ Q: What happens when a peer receives two identical messages?
- ▶ A: Unique identifiers are necessary!

Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

P2P Overlay Networks

Dissemination Algorithms

Security

Consistency

Bitcoin

Conclusion

Disemination Algorithms

Problem: broadcast a message to all the peers.

Solutions:

- ▶ Flooding
- ▶ Gossip

Flooding Algorithms

- ▶ Each peer forwards the new information to its connections.
- ▶ To reduce bandwidth consumption
 - peers can first exchange hashes of messages,
 - and transfer only those unseen before.

Gossip Algorithms

Principle:

Gossiping is the endless process of
 randomly choosing two members
 and subsequently letting these two exchange information
(Kermarrec/Van Steen, "Gossiping in distributed systems")

Gossip Algorithm: Rumor Mongering

1. Peers are initially ignorant.
2. When a peer learns a new message it becomes a hot rumor.
3. Periodically, the chooses a random peer
and sends (pushes) the rumor to it.
4. Eventually, the peer loses interest in spreading the rumor.

Gossip Algorithm: Anti-Entropy

1. Each peer p periodically contacts a random peer q .
2. p and q engage information:
 - updates from p are transferred to q (push),
 - or vice-versa (pull),
 - or in both direction (push-pull).

Rumor Propagation

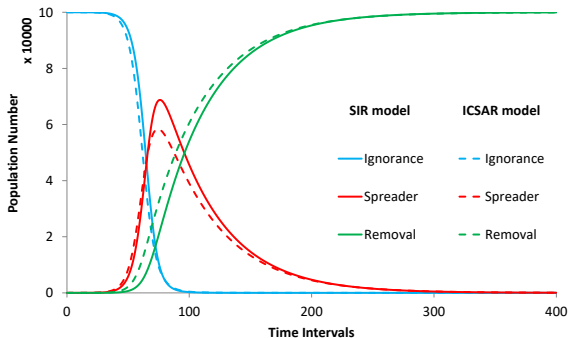


Figure from Zhang & all: Dynamic 8-state ICSAR rumor propagation model considering official rumor refutation.

Peer Sampling in Gossiping Algorithms

- ▶ Each peer periodically exchanges information with a set of peers.
- ▶ The choice of this set is crucial.
- ▶ Ideally, it should be a uniform random sample of all online peers.

Peer Sampling in Gossiping Algorithms

- ▶ Each peer periodically exchanges information with a set of peers.
- ▶ The choice of this set is crucial.
- ▶ Ideally, it should be a uniform random sample of all online peers.

Issue:

- ▶ Infeasible to maintain a list with all peers in the P2P overlay.

Peer Sampling: A Gossip-like Solution

- ▶ Every peer maintains a relatively small list of peers.
- ▶ It periodically modifies it using a gossiping procedure:
 - it selects a peer
 - and the two exchange (swap) parts of their lists

Overview Disemination Algorithms

Problem: broadcast a message to all the peers.

Solutions:

- ▶ Flooding: robust,
less efficient ($O(n^2)$)
- ▶ Gossip: robust,
efficient ($O(n \log n)$),
relative high latency,
only probabilistic guarantees for reliability

Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

P2P Overlay Networks

Disemination Algorithms

Security

Consistency

Bitcoin

Conclusion

Blockchain: Challenges and Approaches

- ▶ Q: How to make sure Alice does not “double spends”?

Blockchain: Challenges and Approaches

- ▶ Q: How to make sure Alice does not “double spends”?
- ▶ Every peer verifies the transaction before adding it to the ledger.

Blockchain: Challenges and Approaches

- ▶ Q: How to make sure Alice does not “double spends”?
- ▶ Every peer verifies the transaction before adding it to the ledger.
- ▶ Q: What can get wrong?

Blockchain: Challenges and Approaches

- ▶ Q: How to make sure Alice does not “double spends”?
- ▶ Every peer verifies the transaction before adding it to the ledger.
- ▶ Q: What can get wrong?
- ▶ Nota bene: decentralized p2p network

Sybil Attack

- ▶ Alice sets up multiple identities

Sybil Attack

- ▶ Alice sets up multiple identities
- ▶ Creating a new peer requires:
 - memory/disk space and
 - processor time.
- ▶ Several peer can co-exist on the same computer.
- ▶ The majority of the peers belong to Alice.

Solution: Proof of Work

Every peer:



1.	Collect and verify submitted transactions to form a block
2.	Proof of Work: Solve a puzzle (competition between peers)
3.	The winner can append its block to the chain Hence, it announces the block to the network

Step “Proof of Work” as a competition

- ▶ Q: Why?



Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

P2P Overlay Networks

Disemination Algorithms

Security

Consistency

Bitcoin

Conclusion

Information Propagation

Issue: Network becomes inconsistent once a new block is generated

Case of study Bitcoin 2013

- ▶ Average time till a node receives a new block : 12.6 seconds
- ▶ Longtail: 5% of nodes do not have the new block after 40seconds
- ▶ Consequence: blockchain forks
- ▶ 169 forks were observed during a period of 10,000 generated blocks

Decker and Wattenhofer IEEE P2P'13

Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

Bitcoin

- Data Structures

- Nakamoto Consensus

- Proof-of-Work

- Consensus

- Issues


Conclusion

Bitcoin

- ▶ The block chain provides Bitcoin's public ledger.
ledger = the principal book recording the commercial transactions

The Blockchain

- ▶ A block is an ordered and timestamped record of transactions.
- ▶ Each full node in the Bitcoin network independently stores a block chain.



Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

Bitcoin

Data Structures

Nakamoto Consensus

Proof-of-Work

Consensus

Issues

Conclusion

Bitcoin: “Nakamoto Consensus”

- ▶ New transactions are broadcasted to all peers (miner).
- ▶ Each miner collects¹ new transactions into a block.
- ▶ In each round a random² miner gets the privilege add its block to the chain. To do so, the miner broadcasts the block in the network.
- ▶ Other miners accept the block if all the transactions in it are valid.
- ▶ Miners express their acceptance of the block by including its hash in the next block they create.

¹Some systems prioritise transactions those fees are higher (a better reward).

²Actually, the winner of the proof-of-work puzzle.

Miners

- ▶ Miners are peers that:
 - receive and validate transactions,
 - group them into a block
- ▶ Miners race to solve a puzzle (poof-of-work) in order to publish their³ block
- ▶ Miners invest a large amount of computer in the race.
- ▶ The reward is valid only if other nodes accept the published block.

³Create/earn a bitcoin reward for each block published (“mined”)

Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

Bitcoin

Data Structures

Nakamoto Consensus

Proof-of-Work

Consensus

Issues

Conclusion


Proof-of-Work: The Hash Lottery

- ▶ Difficult Problem (Proof of Work) for the miner:
finding a number (known as a nonce) that
when nonce is added to a given input data,
the value of the hash begins with a number of zeros:
 $H(\text{nonce} \mid \text{data}) = 0000000\text{xxxx}$
- ▶ In the Bitcoin world this is what “mining” is.
- ▶ A miner’s win probability is proportional to its compute power.
In this way, the next miner to publish is selected randomly.

Why is this problem important?

Proof-of-Work in Bitcoin:

- ▶ Effect: a lot of hash-power is spent on guessing winning lottery numbers that satisfy the difficulty of the problem.
- ▶ Incentive: "Guess" numbers in order to obtain the reward from the network.



As of January 2018, the global hashrate is about 17 exahash per second (EH/s).

The Payout

- ▶ The node that finds the best solution to the challenge is granted a reward.
- ▶ Originally in Bitcoin it was 50 new coins.
- ▶ If a nodes receives two competing solutions it selects the one that has:
 - higher number of transactions included in the block,
 - higher number of zeros.

Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

Bitcoin

Data Structures

Nakamoto Consensus

Proof-of-Work


Consensus

Issues

Conclusion

Arriving at Consensus

How to avoid forks in the block chain?



- ▶ When a miner links to a block, it accepts as the head of the valid chain.
- ▶ If others disagree, then the miner's block is worthless.
... so miners have an incentive to get it right!
- ▶ If a block A is posted “too late” (e.g., the block in grey) or is invalid, other miners ignore it, and build the chain in another direction.
- ▶ The longest chain wins: defines the “next puzzle”

Transaction Confirmation

- ▶ Transaction accepted in a candidate block \Rightarrow the transaction is valid.
- ▶ Confirmation: Every new block accepted into the chain after the block containing the transaction counts as a confirmation.
- ▶ Coins are not considered mature until there have been 6 confirmations.
means 1h if a 10 minute block cadence
- ▶ New mined coins are not valid until about 120 confirmations.
assures that a node with more than 51% of the total hash-power
does not pull off fraudulent transactions
- ▶ Checkpoints to insure that every to prevent complete history rollback.

Agenda

Crypto Principles (from Eduard Felten's tutorial)

Distributed Architecture

Bitcoin

- Data Structures

- Nakamoto Consensus

- Proof-of-Work

- Consensus

- Issues

Conclusion

Node with Limited Functionality

Leave out some functionality

- ▶ Calculation of proofs of work (mining)
- ▶ Store only block headers instead of complete blocks
 - incomplete verification of transactions
 - trust others (verification of the transaction + proof of work)
- ▶ Acceptance of incoming connections

Node with Limited Functionality

Leave out some functionality

- ▶ Calculation of proofs of work (mining)
- ▶ Store only block headers instead of complete blocks
 - incomplete verification of transactions
 - trust others (verification of the transaction + proof of work)
- ▶ Acceptance of incoming connections

Potential attacks on SPV

- ▶ Attacker controls victim's internet connection.
- ▶ It double-spends and computes own proof(s) of work.
- ▶ Victim does not check for double-spending.

Privacy issues

Transaction graph

- ▶ Link different Bitcoin addresses of a user

Peer-to-Peer Network

- ▶ Find origin (IP address) that broadcasted a transaction:
 - create fake identities and let many peers, and connect to you
 - connect to peers that accept incoming connections
 - as described in a talk by Dan Kaminsky, CC Congress 2011

Scalability

Broadcast

- ▶ If Bitcoin would process VISA's 2,000 transactions per second,
→ this would lead to a traffic of 8 MB/s

Denial-of-Service Attacks

- ▶ Goal: overload or stop the transmission of transactions/blocks
- ▶ Hard to counteract in a P2P overlay network
- ▶ Limit block size

Conclusion Distributed Consensus

Distributed consensus can allow:

- ▶ Distrustful parties to maintain clean state
- ▶ Completely unambiguous rules about validity
- ▶ Removing authentication and identity as essential

Other applications to Blockchain Technology

- ▶ Registeries
- ▶ Authoritative Systems of Record
- ▶ Directory Services
- ▶ Timestamping Services (“Proof of Existence”)
- ▶ Counter-party Exchanges